

CSCI 4061: Introduction to Operating Systems

Project 3: Multi-Threaded Web Server

Due: Dec 10th at 11:55pm.

1. Overview

The purpose of this lab is to construct a multi-threaded web server using POSIX threads (pthreads) in C language to learn about thread programming and synchronization methods. Your web server should be able to handle any file type: HTML, GIF, JPEG, TXT, etc. and of any arbitrary size. It should handle a limited portion of the HTTP web protocol (namely, the GET command to fetch a web page / files). We will provide pieces of code (some already compiled into object files, and some source) that will help you complete the lab.

2. Description

Your server will be composed of two types of threads: **dispatcher threads** and **worker threads**. The purpose of the dispatcher threads is to repeatedly accept an incoming connection, read the client request from the connection, and place the request in a queue. We will assume that there will only be one request per incoming connection. The purpose of the worker threads is to monitor the request queue, retrieve requests and serve the request's result back to the client. The request queue is a bounded buffer and will need to be properly synchronized (using CVs).

You will use the following functions which have been precompiled into object files that we have provided (full documentation of these functions has been provided in util.h. More will be said below about how to use these functions). Since the server will block, you will need to ^C it to terminate.

You can assume that all the below functions are thread safe.

```
void init (int port); // run ONCE in your main thread
```

Each of the next two functions will be used in order, in a loop, in the dispatch threads:

```
int accept_connection (void);  
int get_request (int fd, char *filename);
```

These functions will be used by the worker threads after handling the request

```
int return_result (int fd, char *content_type, char *buf, int  
numbytes);  
int return_error(int fd, char *buf); // for error.
```

3. Thread pool

Your server should create a static/fixed pool of dispatcher and worker threads when the server starts. The dispatcher thread pool size should be `num_dispatch` and the worker thread pool size should be `num_workers` (See section 8 for details). You can decide whether to create joinable threads or detachable threads.

Extra Credit:

You can implement a variation by dynamically varying worker thread pool size instead of having a fixed number of workers. You can run one of the above two variations based on the `dynamic_flag` (See section 8 for details). We expect you to be creative and come up with your own policy to dynamically increase/decrease the pool size depending on the server load. But note that, there needs to be at-least one worker thread at any point of time. You must explain the policy you implemented in the ReadMe file.

Hints:-

1. You can think in terms of finding a relation between pending requests and number of worker threads. This is just one way of thinking. You can come up with your own creative policy!
2. You can keep a separate daemon thread for implementing your policy.

4. Incoming Requests

An HTTP request has the form: `GET /dir1/subdir1/.../target_file HTTP/1.1` where `/dir1/subdir1/.../` is assumed to be located under your web tree root location. Our `get_request()` automatically parses this for you and gives you the `/dir1/subdir1/.../target` file portion and stores it in the `filename` parameter. Your web tree will be rooted at a specific location, specified by one of the arguments to your server (See section 8 for details). For example, if your web tree is rooted at `/home/user/joe/html`, then a request for `/index.html` would map to `/home/user/joe/html/index.html`. You can `chdir` into the Web root to retrieve files using relative paths.

5. Returning Results

You will use `return_result()` to send the data back to the web browser from the worker threads provided the file was opened and read correctly, or the data was found in the cache. If there was any problem with accessing the file, then you should use `return_error()` instead. Our code will automatically append HTTP headers around the data before sending it back to the browser. Part of returning the result is sending back a special parameter to the browser: the content-type of the data. You may make assumptions based on the extension of the files as to which content-type they are:-

- Files ending in .html or .htm are content-type “text/html”
- Files ending in .jpg are content-type “image/jpeg”
- Files ending in .gif are content-type “image/gif”
- Any other files may be considered, by default, to be of content-type “text/plain”.

6. Caching

To improve runtime performance, you need to implement caching which stores cache entries in memory for faster access. When a worker serves a request, it will look up the request in the cache first. If the request is in the cache (Cache HIT), it will get the result from the cache and return it to the user. If the request is not in the cache (Cache MISS), it will get the result from disk as usual, put the entry in the cache and then return result to the user. The cache size can be defined by an argument and you need to log information about the cache (HIT or MISS) with time (see section 7,8 for more details). How to implement caching is totally up to you. You can implement a simple cache replacement policy like random policy or FIFO policy to choose an entry to evict when the cache is full.

(**Extra Credit:** Implement LRU or LFU cache replacement policy.)

7. Request Logging

From the worker threads, you must carefully log each request (normal or error-related) to a file called “web_server_log” and also to the terminal (stdout) in the format below. You must also protect the log file from multiple threads attempting to access it simultaneously.

[threadId][reqNum][fd][Request string][bytes/error][time][Cache HIT/MISS]

- **threadId** is an integer from 0 to num_workers -1 indicating thread ID of request handling worker.
- **reqNum** is total number of requests a specific worker thread has handled so far, including current request.
- **fd** is the file descriptor given to you by accept_connection() for this request
- **Request string** is the filename buffer filled in by the get request function
- **bytes/error** is either the number of bytes returned by a successful request, or the error string returned by return error if an error occurred.
- **time** is the service time of a request

- **Cache HIT/MISS** is either “HIT” or “MISS” depending on whether this specific request was found in the cache or not.

The log (in the “web_server_log” file and in the terminal) should look something like the example below.

```
[8][1][5][image/jpg/30.jpg][17772][200ms][MISS]
[9][1][5][image/jpg/30.jpg][17772][3ms][HIT]
```

8. How to run the server

Your server should be run as:

```
% ./web_server port path num_dispatch num_workers dynamic_flag qlen cache_entries
```

The server will be configurable in the following ways:

- `port` number can be specified (you may only use ports 1025 - 65535 by default)
- `path` is the path to your web root location from where all the files will be served
- `num_dispatcher` is how many dispatcher threads to start up
- `num_workers` is how many worker threads to start up
- `dynamic_flag` indicates whether the worker thread pool size should be static or dynamic. By default, it should be 0.
- `qlen` is the fixed, bounded length of the request queue
- `cache_entries` is the number of entries available in the cache (an alternative way to do this would be to specify the maximum memory used by the cache, but we are just using a limit on the number of entries)

9. Provided Files and How to Use Them

We have provided many functions which you must use in order to complete this assignment. We have handled all of the networking system calls for you. We have also handled the HTTP protocol parsing for you. Some of the library function calls assumes that the program has “chdir”ed to the web tree root directory. You need to make sure that you chdir to the web tree root somewhere in the beginning.

We have provided a makefile you can use to compile your server. Here is a list of the files we have provided.

1. `server.c` : You only need to modify this file to implement a server.

2. `makefile` : You can use this to compile your program using our object files, or you can make your own. You can study this to see how it compiles our object code along with your server code to produce the correct binary executables.
3. `util.h` : This contains a very detailed documentation of each function that you must study and understand before using the functions.
4. `util.o` : This is the compiled code of the functions described in `util.h` to be used for the web server. Compile this into your multi-threaded server code and it will produce a fully-functioning web server.
5. `testing.tar.gz` : This file includes images, texts and url files to test your server. See section 10 for more detailed information.

10. How to test your server

We will test your server with “`wget`”, the non-interactive network downloader. After you run the server, open a new terminal to test the server. You can try to download a file using this command:

```
-> wget http://127.0.0.1:9000/image/jpg/29.jpg
```

(Please note that 127.0.0.1 means localhost)

This command will try to download the file at `root/image/jpg/29.jpg`. If it failed to download the file for some reason, it will show an error message. You can also test your server with any web browser (Internet Explorer, Chrome, Firefox, and so on). We will provide “`testing.tar.gz`” to make testing the server easier. Once you extract it, you can find an instruction file “`how_to_test`” which explains how to use and test your program with these files. The server should be able to serve requests from other machines, so you need to run the server on a machine and run “`wget`” command on another machine to test your server (use the server's IP address instead of 127.0.0.1). Note, if you try to connect to the server which runs on CSELabs machine with your own machine (laptop), it may not be able to connect to the server because CSELabs machines are protected by a firewall.

Testing Concurrency: Bash has a nifty command called `xargs`, which allows a set of arguments to be piped to a command such that multiple executions will be run concurrently.

The format of the command is '`xargs -n num_args -P num_procs cmd`'

So, for our purposes, the command “`cat urls | xargs -n 1 -P 8 wget`” will run `wget` 8 times simultaneously (`-P 8`) with 1 argument each time (`-n 1`) from the pipe produced by `cat urls`.

NOTE: If -P is given an argument that is smaller than the number of args in the pipe, then multiple sets of size P will be run, with each set being concurrent.

Ex. If I had 10 arguments in the pipe, and I set -P 5, 2 sets of 5 concurrent processes will run. This can be used to test that your server really does permit concurrent execution of multiple requests.

11. Simplifying Assumptions

- The maximum number of dispatcher threads will be 100.
- The maximum number of worker threads will be 100.
- The maximum length of the request queue will be 100 requests.
- The maximum size of the cache will be 100 entries.
- Any HTTP request for a filename containing two consecutive periods or two consecutive slashes (“..” or “//”) will automatically be detected as a bad request by our compiled code for security.

12. Documentation

You must include a README containing the following information:

- Your group ID and each member’s name and X500.
 - How to compile and run your program.
 - A brief explanation on how your program works.
 - Explanation of caching mechanism used.
 - Explanation of your policy to dynamically change the worker thread pool size.
 - Contributions of each team member towards the project development.
 - Within your code you should use one or two sentences to describe each function that you wrote.
- You do not need to comment every line of your code. However, you might want to comment portions of your code to increase readability.

At the top of your README file, please include the following comment:

```
/* CSci4061 Fall 2018 Project 3
```

```
* Name: <full name1>, <full name2>
```

```
* X500: <X500 for first name>, <X500 for second name> */
```

13. Deliverables

1. Files containing your code
2. A README file (readme and C code should indicate this is project 3). README should contain information on every group member's individual contributions. All files should be compressed into a single tar file, named `Project_3_group_<your_group_number>.tar` and submitted through canvas. This is your official submission that we will grade. We will only grade the most recent and on time submission. Failing to follow the submission instructions may result in a 5% deduction.
3. Performance analysis report
 - a. Run a few commands which will result in a mix of cache HIT and cache MISS. And explain the results based on the observed service time for those requests.
 - b. Run your server with 100 dispatcher threads and different number of worker threads (1, 2, 5, 10, 20, 50, 75, 100).
 - i. Run the xargs test with **bigurls** file.
 - ii. Plot the time taken for all the requests to complete. (You can use the time command in linux and report the real time)
 - iii. Explain the observed results based on theory and your implementation.

14. Grading (Tentative)

10% README, Documentation within code, coding and style (indentation, readability of code, use of defined constants rather than numbers).

10% Performance Analysis

80% Correctness, error handling, meeting the specification. Broken down as follows:

- 10% for handling one request successfully.
- 10% for handling multiple requests.
- 30% for handling multiple concurrent requests (test with xargs).
- 30% for implementing caching.

Again, this must all be done with error handling and according to the spec for full points.

Grading will be done on the CSELabs machines. All project policies enforced on other projects should be assumed to apply unless stated hereafter on the canvas site.